



Machine Learning

Lecture 6

Autoencoders




Deep Autoencoders




Latent Representation


Sampling Distribution



Latent Representation




Autoencoders




Better Alternative

Sampling Distribution




Better Alternative

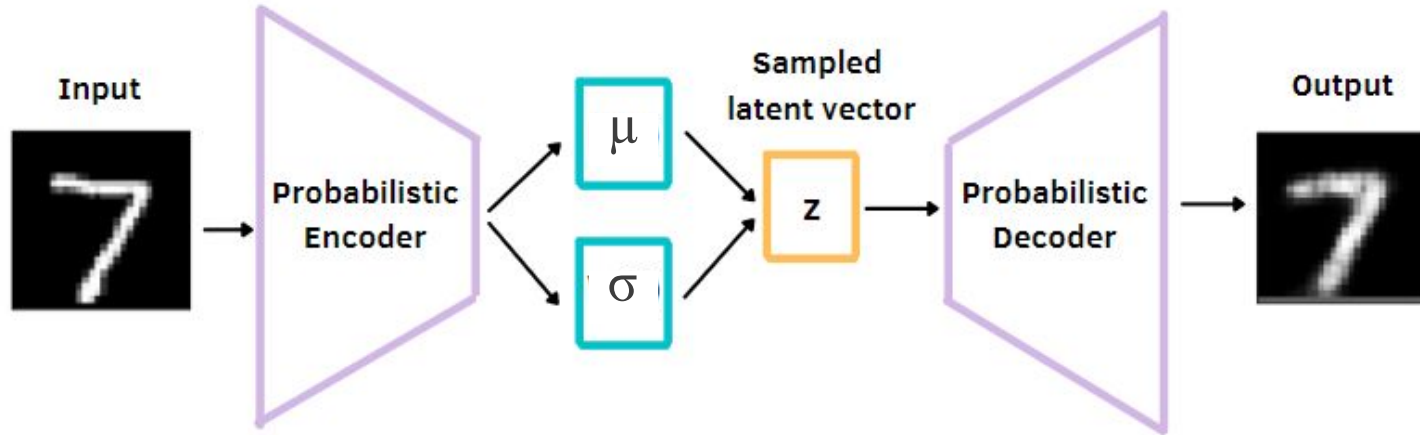
Sampling Distribution



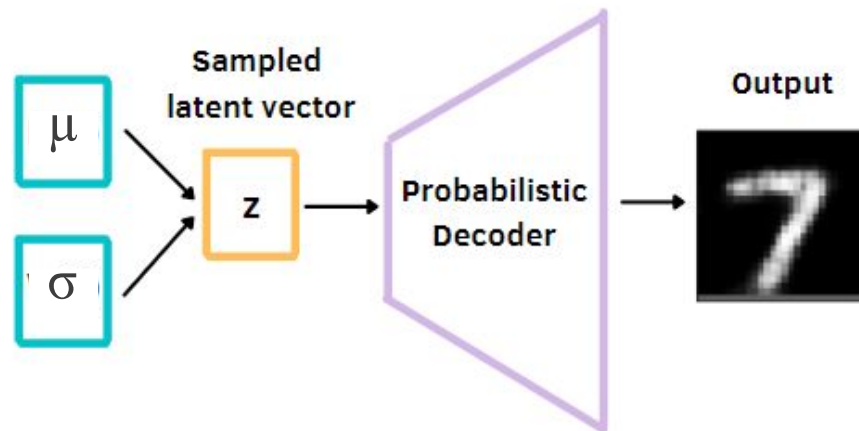
Autoencoders



Variational Autoencoders (VAEs)



Generating with VAEs

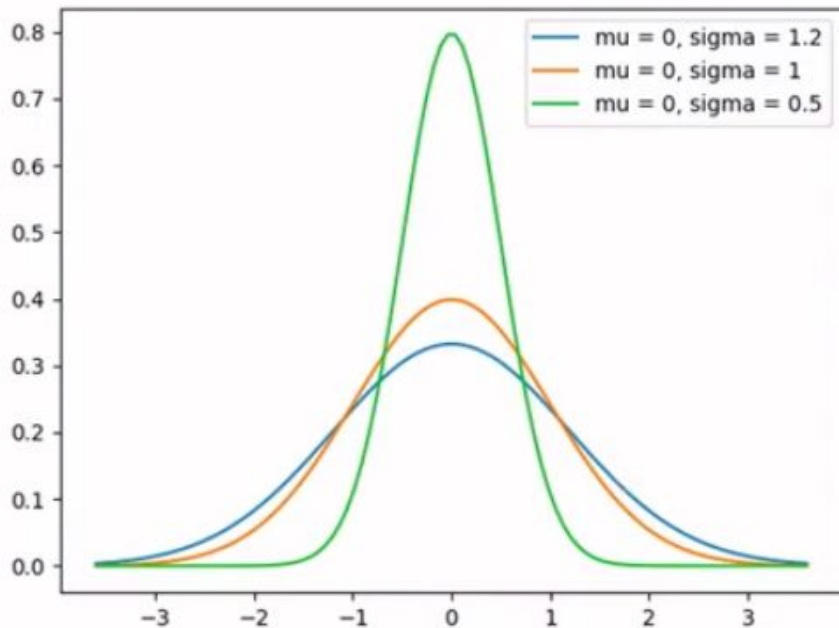


Few Definitions

- Gaussian Distribution
 - $z \sim N(\mu, \sigma)$.
 - μ - mean
 - σ - standard deviation


Few Definitions

- Gaussian Distribution
 - $z \sim N(\mu, \sigma)$.
 - μ - mean
 - σ - standard deviation



Few Definitions


- Gaussian Distribution
 - $z \sim N(\mu, \sigma)$.
 - μ - mean
 - σ - standard deviation





Reparameterization Trick

- We want to use gradient descent to learn the model's parameters.
- Given z drawn from $q_{\theta}(z|x)$, how do we take derivatives of (a function of) z w.r.t. θ ?
- We can reparameterize: $z = \mu + \sigma \odot \epsilon$
- $\epsilon \sim N(0, 1)$
- Now we can take derivatives of (functions of) z w.r.t. μ and σ .
- Output of $q_{\theta}(z|x)$ is vector of μ 's and vector of σ 's.


Reparameterization Trick



If (μ, σ)
are known






Reparameterization Trick



If (μ, σ) are known,

$$g_{\mu, \sigma}(\epsilon) = \mu + \sigma \odot \epsilon$$
$$\epsilon \sim N(0, 1)$$
$$z = g_{\mu, \sigma}(\epsilon)$$

Reparameterization Trick



Loss Functions

Reconstruction Loss → **Binary Cross Entropy Loss**



Input



Output

$$-\frac{1}{N} \sum_{i=1}^N y_i \cdot \log(p(y_i)) + (1 - y_i) \cdot \log(1 - p(y_i))$$

Latent Loss → **KL Divergence Loss**

$$D_{KL}(P||Q) = \int P(x) \log \frac{P(x)}{Q(x)} dx$$

