


Machine Learning


Lesson 2

Machine Learning


Unsupervised




Supervised



Reinforcement




Supervised Learning



Regression

Classification


Supervised Learning



X	Y
1	2
3	4
3	5
6	7
7	7
3	4

.
. .
. .

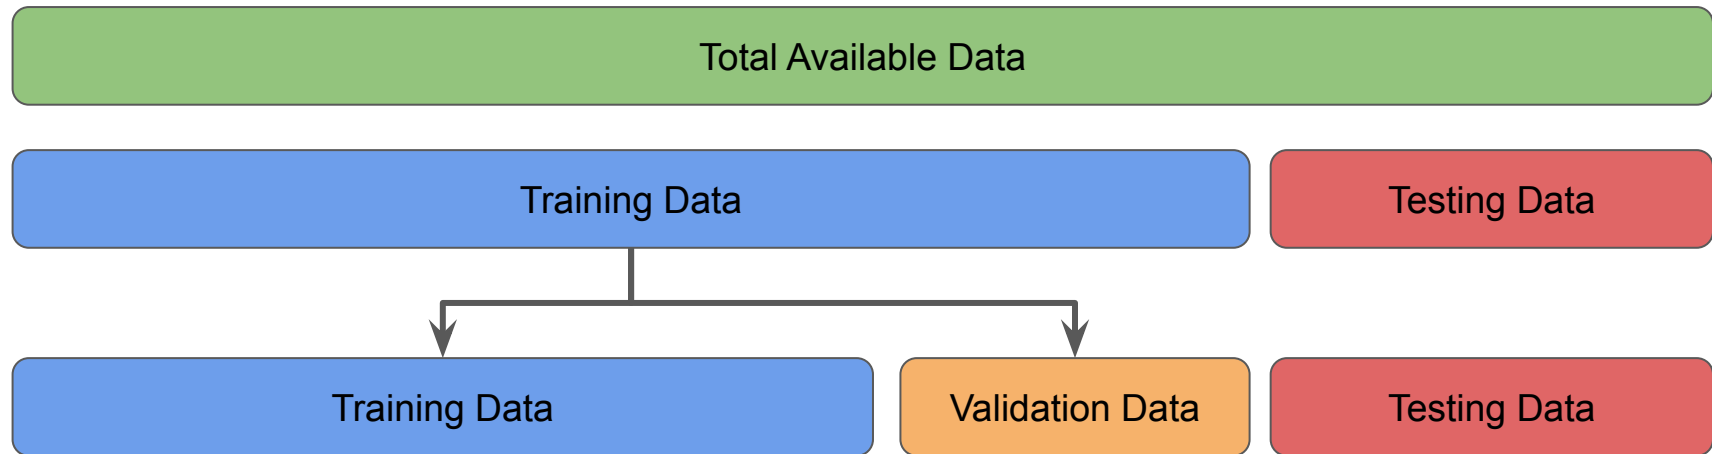
Supervised Learning



x1	x2	Label
1	2	O
3	4	O
3	2	X
6	3	X
7	4	X
3	1	O

⋮ ⋮ ⋮
⋮ ⋮ ⋮
⋮ ⋮ ⋮




Dataset




Supervised Machine Learning Algorithms

- Nearest Neighbor
- Naive Bayes
- Decision Trees
- Linear Regression
- Logistic Regression
- Support Vector Machines (SVM)
- Neural Networks

Linear Regression



Linear Regression



$$y = \theta x + b$$

$$MSE = \frac{1}{N} \sum_{i=1}^N (\hat{y}_i - y)^2$$

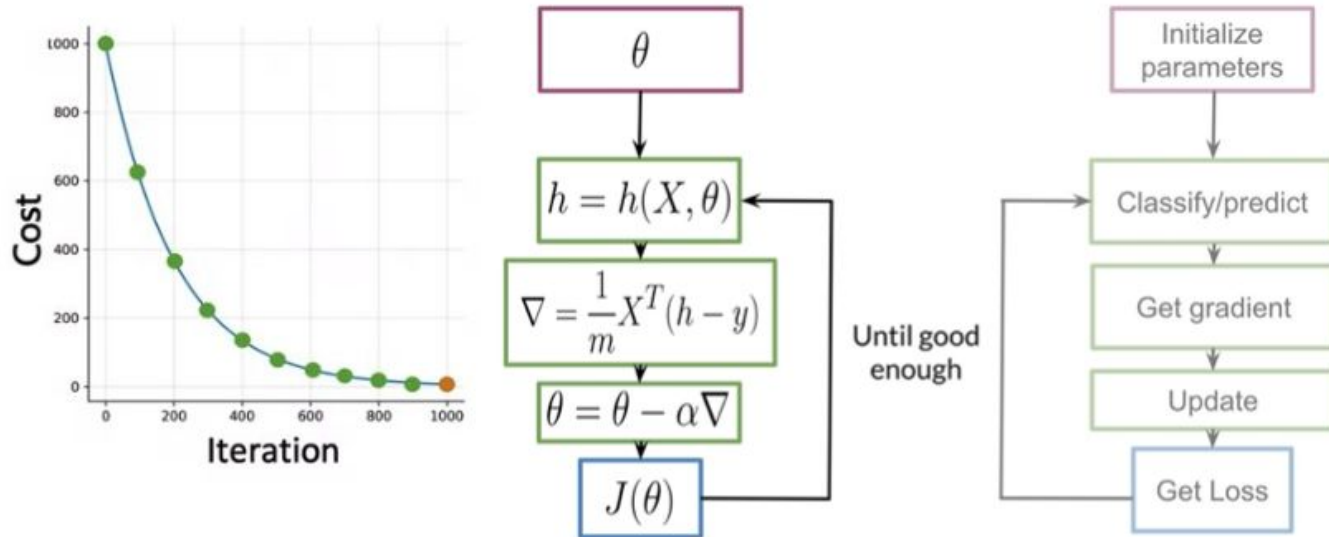
$$L(\theta, b) = \frac{1}{N} ((x\theta + b) - y)^2$$

$$\frac{dL}{d\theta} = \frac{1}{N} x^T ((x\theta + b) - y)$$

$$\frac{dL}{db} = \frac{1}{N} ((x\theta + b) - y)$$

Linear Regression: Training

Usually you keep training until the cost converges. If you were to plot the number of iterations versus the cost, you should see something like this:



You initialize your parameter θ , that you can use in your equation, you then compute the gradient that you will use to update θ , and then calculate the cost. You keep doing so until good enough.


Linear Regression : Gradient Descent

$$\min_{\theta} J(\theta)$$

$$\theta \leftarrow \theta - \alpha \frac{d}{d\theta} J(\theta)$$




$$\frac{d}{d\theta} J(\theta) = (h(x, \theta) - y) x$$

$$\theta \leftarrow \theta - \alpha (h(x, \theta) - y) x$$






$\alpha \rightarrow$ learning rate

Logistic Regression



Logistic Regression



$$x_2 = \theta \cdot x_1 + b$$


Logistic Regression

But we don't need \mathbf{x}_2 , we need probability that a point is on either side of that line.

So we use the sigmoid function,

$$f(x) = \frac{1}{1 + e^{-x}}$$

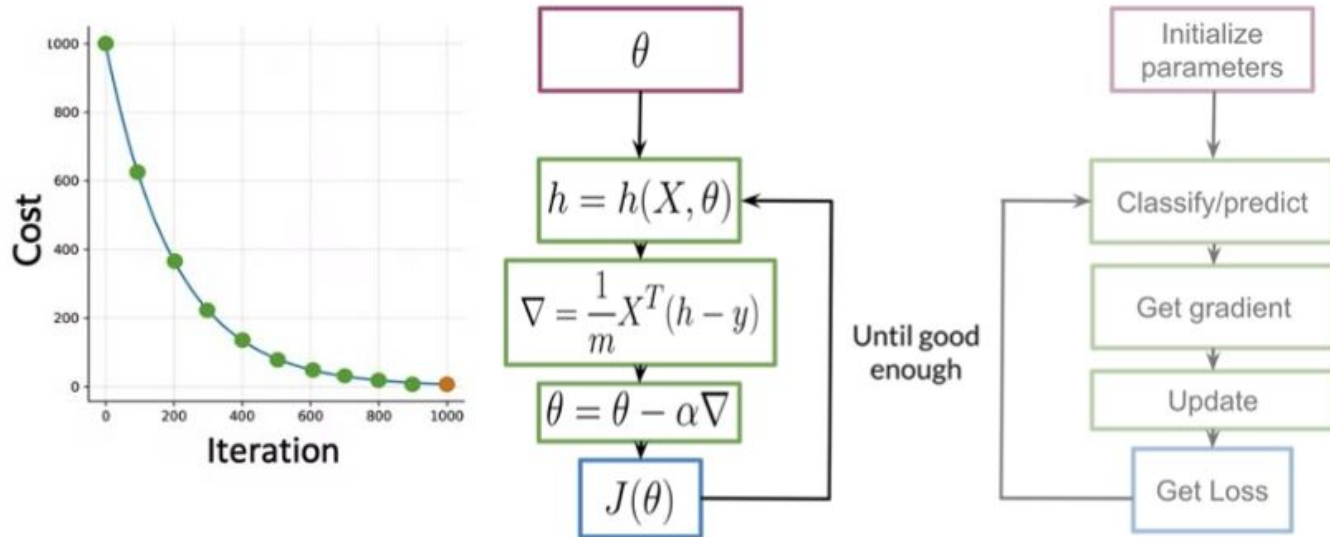
$$P = f(\theta x_1 + b) = \frac{1}{1 + e^{-(\theta x_1 + b)}}$$



Note that as $(\theta x_1 + b)$ gets closer and closer to $-\infty$ the denominator of the sigmoid function gets larger and larger and as a result, the sigmoid gets closer to 0. On the other hand, as $(\theta x_1 + b)$ gets closer and closer to $+\infty$ the denominator of the sigmoid function gets closer to 1 and as a result the sigmoid also gets closer to 1.

Logistic Regression : Training


Usually you keep training until the cost converges. If you were to plot the number of iterations versus the cost, you should see something like this:



You initialize your parameter θ , that you can use in your sigmoid, you then compute the gradient that you will use to update θ , and then calculate the cost. You keep doing so until good enough.

Logistic Regression : Cost Function

$$J(\theta) = - [y \log h(x, \theta) + (1 - y) \log (1 - h(x, \theta))]$$



As you can see in the picture above, if $y = 1$ and you predict something close to 0, you get a cost close to ∞ . The same applies for then $y=0$ and you predict something close to 1. On the other hand if you get a prediction equal to the label, you get a cost of 0. In either, case you are trying to minimize $\mathbf{J}(\theta)$.


Logistic Regression : Gradient Descent

$$\min_{\theta} J(\theta)$$

$$\theta \leftarrow \theta - \alpha \frac{d}{d\theta} J(\theta)$$

$$\frac{d}{d\theta} J(\theta) = (h(x, \theta) - y) x$$

$$\theta \leftarrow \theta - \alpha (h(x, \theta) - y) x$$



$\alpha \rightarrow$ learning rate